



ML Lab Week 10 SVM Lab Instructions

SRN : PES2UG23CS366

SECTION : F

SEM – 5

1. Objective

The goal of this lab is to understand and implement Support Vector Machine (SVM) classifiers. You will train SVMs using three different kernels: **Linear**, **Radial Basis Function (RBF)**, and **Polynomial**, on distinct datasets. You will then evaluate their performance using standard classification metrics and visualize their decision boundaries to see how they separate data.

2. Core Concepts

- **Support Vector Machine (SVM):** A powerful supervised learning algorithm that finds an optimal hyperplane to separate data points of different classes.
- **Kernel Trick:** A technique that allows SVMs to solve non-linear problems by transforming data into a higher-dimensional space.
 - **Linear Kernel:** Creates a straight-line decision boundary.
 - **RBF Kernel:** Creates a complex, non-linear boundary, like a circle or a wave.
 - **Polynomial Kernel:** Creates a curved, polynomial decision boundary.
- **Hard vs. Soft Margin:** The parameter C in SVMs controls the trade-off between maximizing the margin and minimizing the classification error. A large C leads to a hard margin (less tolerance for misclassification), while a small C leads to a soft margin (more tolerance).

3. Deliverables

1. .


- ♦ **Content Requirements:**

1. **Screenshots Provide clearly labeled screenshots for all the results generated by your notebook. You must include a total of 14 screenshots, divided as follows:**

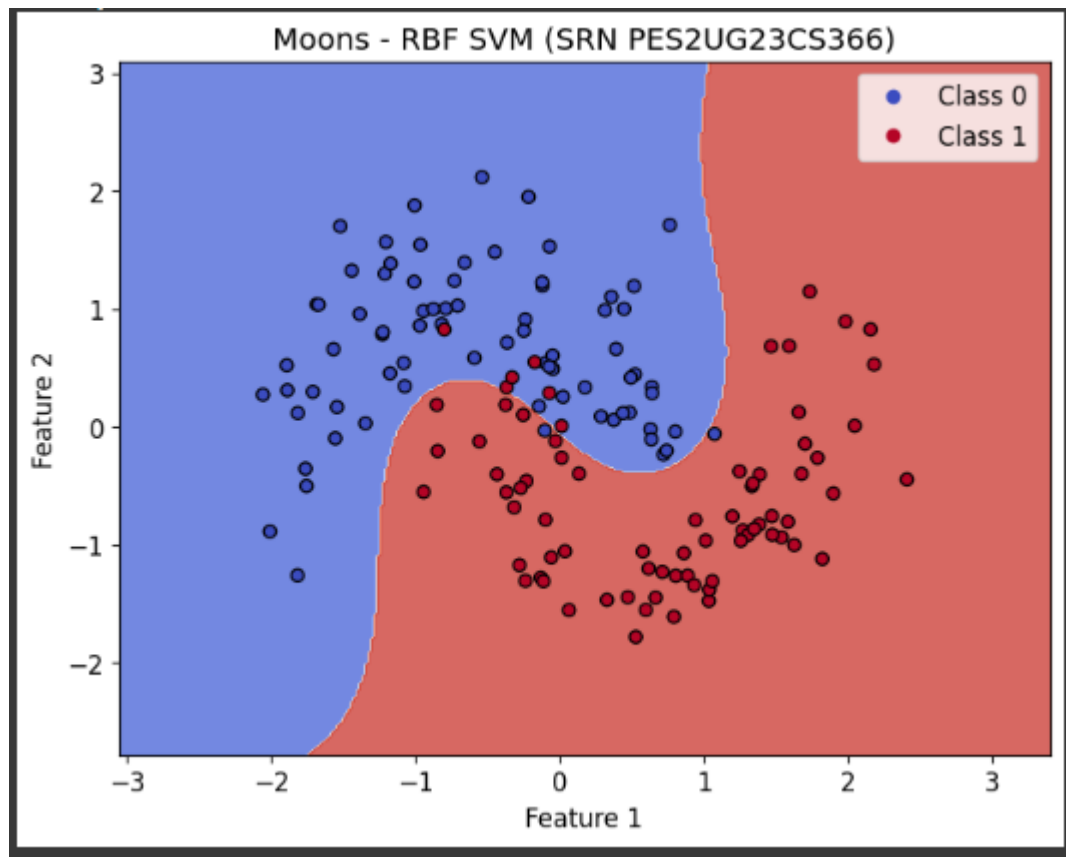
- **Training Results (6 Screenshots):** Capture the classification report output for each model.

- **Moons Dataset (3 screenshots):**


1. Classification Report for SVM with LINEAR Kernel with SRN



2. Classification Report for SVM with RBF Kernel with SRN




3. Classification Report for SVM with POLY Kernel with SRN

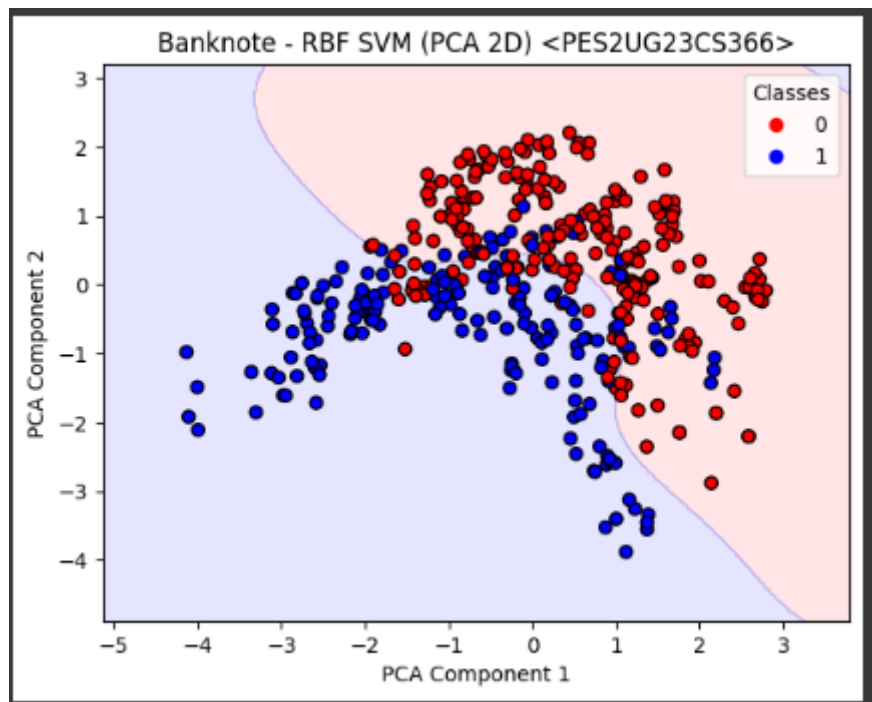


- **Banknote Dataset (3 screenshots):**

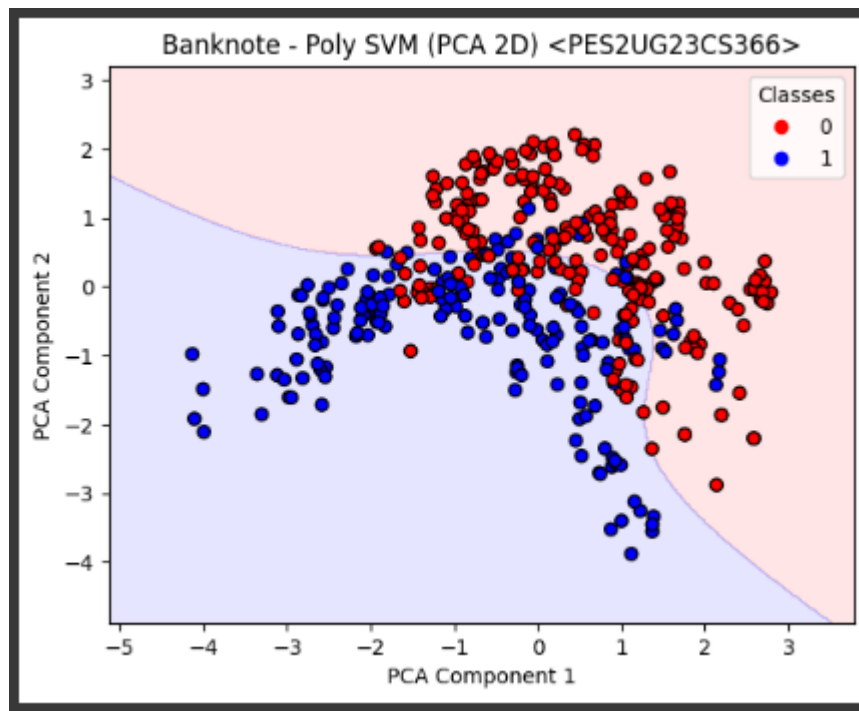
4. Classification Report for SVM with LINEAR Kernel



5. Classification Report for SVM with RBF Kernel



6. Classification Report for SVM with POLY Kernel




- **Decision Boundary Visualizations (8 Screenshots):** Capture the plot for each model's decision boundary.

- **Moons Dataset (3 plots):**


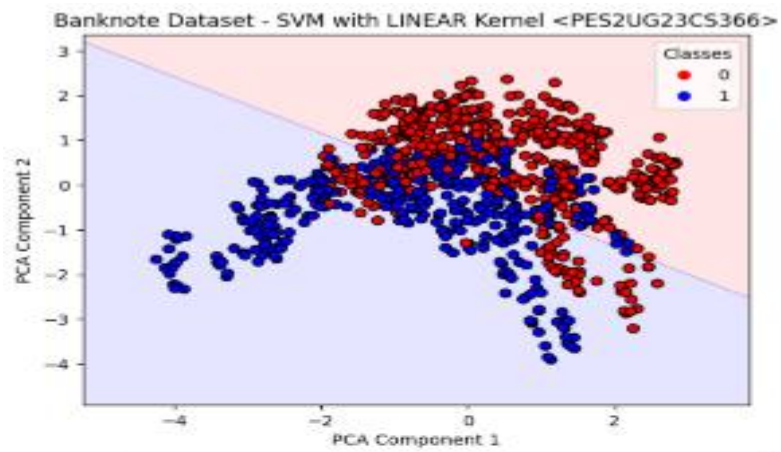
7. Moons Dataset - SVM with LINEAR Kernel
8. Moons Dataset - SVM with RBF Kernel
9. Moons Dataset - SVM with POLY Kernel

Moons Dataset - SVM Kernel Comparison (Decision Boundaries)

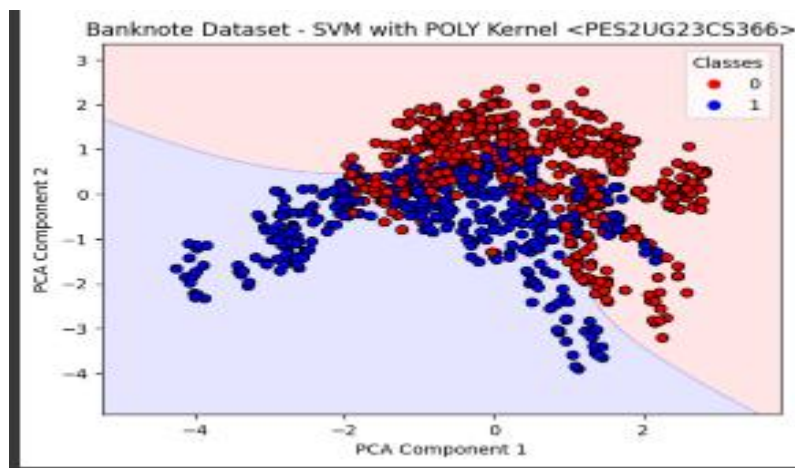


- **Banknote Dataset (3 plots):**

10. Banknote Dataset - SVM with LINEAR Kernel
11. Banknote Dataset - SVM with RBF Kernel


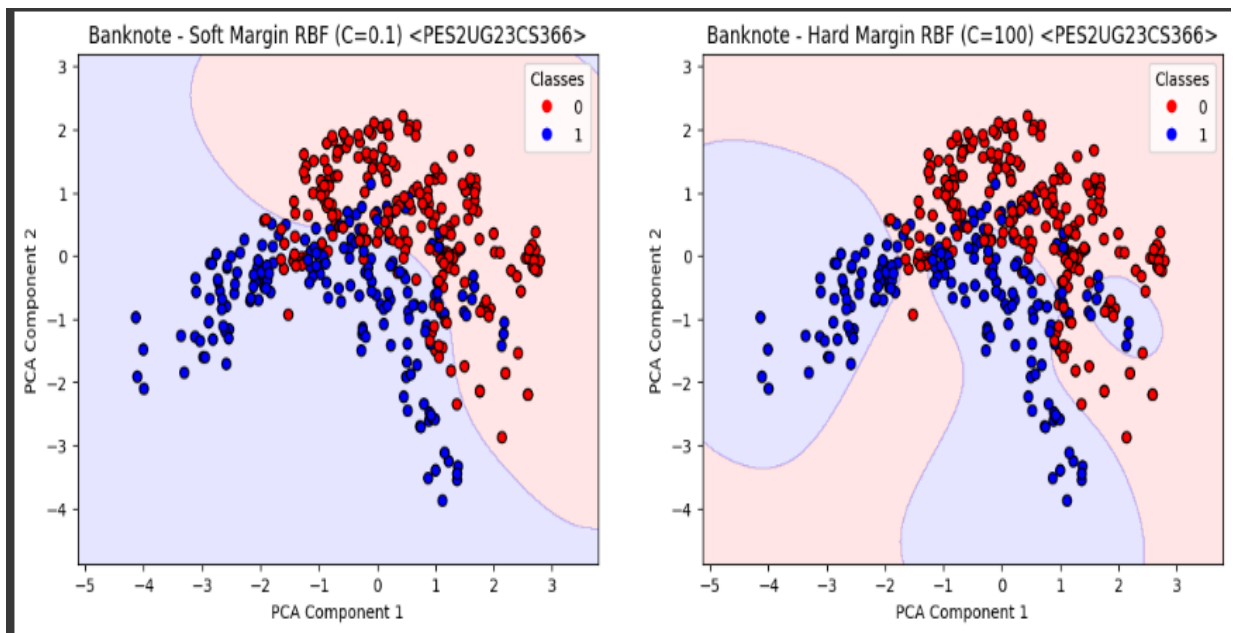



12. Banknote Dataset - SVM with POLY Kernel



● Margin Analysis (2 plots):

- Soft Margin SVM ($C=0.1$)
- Hard Margin SVM ($C=100$)





Analysis Questions:

Analysis Questions for Moons:

1. Based on the metrics and the visualizations, what inferences about the performance of the Linear Kernel can you draw?

Answer :-

Linear kernel — inferences from metrics & visualization

- **Behavior:** Linear SVM produces a straight-line decision boundary.
- **Observed effect:** It **underfits** the Moons data (which is intrinsically non-linear).
- **Evidence:** Lower accuracy (~ 0.84) and lower F1-scores for both classes; many points along the curved moon arcs are misclassified.
- **Conclusion:** Linear kernel is too simple (high bias) for this dataset — it cannot capture the curved class boundary.

2. Compare the decision boundaries of the RBF and Polynomial kernels. Which one seems to capture the shape of the data more naturally?

- **RBF:** Produces *smooth, radial* decision regions that closely follow the curved moon shapes; tends to give the best accuracy (in your run ≈ 0.94 – 0.95).

- **Polynomial (degree 3):** Can fit curved boundaries but may introduce extra wiggles or irregularities (depending on degree and coef0); performance is usually slightly worse than RBF for this dataset.
- **Verdict: RBF** captures the moons more naturally — smoother boundary and better empirical performance.

Analysis Questions for Banknote:

1. In this case, which kernel appears to be the most effective?

- **Typical result:** The **linear kernel** often performs very well on the Banknote Authentication dataset because its features (variance, skewness, kurtosis, entropy) are often linearly separable after scaling.
- **If your experiments show otherwise:** If RBF slightly outperforms linear, it means a small nonlinearity exists; otherwise choose **linear** for simplicity and interpretability.

2. The Polynomial kernel shows lower performance here compared to the Moons dataset. What might be the reason for this?

• Reasons:

- The Banknote dataset is **closer to linearly separable**; an unnecessarily complex polynomial boundary can *overfit* noise instead of improving class separation.
- Polynomial kernels introduce **higher model complexity** (extra curvature / interactions) which is unnecessary when the classes are already separable in the original feature space.
- Sensitive to **feature scaling** and hyperparameters (degree, coef0) — if not tuned, polynomial can underperform.

□ **Conclusion:** Polynomial underperforms because it adds complexity that isn't needed and can fit noise or cause unstable boundaries.

General Margin / C Analysis (Soft vs Hard)

1) Which model produces a wider margin: $C=0.1$ (soft) or $C=100$ (hard)?

- **Wider margin: Soft margin ($C = 0.1$)** produces a wider margin.
 - Low **C** emphasizes maximizing margin over classifying every training point correctly → larger margin, more slack allowed.

2) Why does SVM allow points inside/on the wrong side of the margin in Soft Margin?

- **Reason:** Soft-margin SVM introduces **slack variables** (ξ_i) and a penalty controlled by **C**.
- **Trade-off:** The model trades some training errors (points inside or on the wrong side) for a **larger margin** that usually generalizes better.
- **Primary goal: Maximize the margin** while controlling misclassification via the penalty term — i.e., **good generalization**, not perfect training accuracy.

3) Which model is more likely to overfit?

- **More likely to overfit: Hard margin / high C (C = 100).**
 - High C strongly penalizes misclassification, forcing the classifier to fit training points tightly; this can fit noise and reduce generalization.

4) Which model to trust on a new unseen point? Which C to prefer when data is noisy?

- **Trust more: Soft-margin model (low C)** — more robust to noise and less likely to be misled by outliers.
- **Real-world recommendation:** Start with a **lower C** (e.g., 0.01–1 depending on scale) when data is noisy, then tune C with cross-validation. Low C favors simpler decision boundaries and typically generalizes better; increase C only if validation shows underfitting.

4 . Imagine you receive a new, unseen data point. Which model do you trust more to classify it correctly? Why? In a real-world scenario where data is often noisy, which value of C (low or high) would you generally prefer to start with?

I would trust the **Soft Margin SVM (with low C)** more to classify a new, unseen data point correctly.

Reasoning:

- A **low C value** means the model allows some misclassifications on the training data in exchange for a **wider margin**.
- This wider margin makes the model **less sensitive to noise and outliers**, resulting in **better generalization** to unseen data.
- In contrast, a **high C (hard margin)** model tries to perfectly classify all training samples, which can lead to **overfitting** — it performs well on the training set but poorly on new data.

In a real-world scenario:

- Data is almost always **noisy or imperfect**.
 - Therefore, it is better to **start with a lower C value** (e.g., 0.1 or 1) and later tune it using cross-validation if needed.
 - A smaller C produces a **simpler, smoother decision boundary** that generalizes better.
-

